

# A short talk on parallelizing COSA

Jiading Gai

Center for Research Computing  
University of Notre Dame  
Notre Dame, IN 46556  
jgai@nd.edu

Aug 23, 2010

# Clustering on different subsets of the attributes (COSA):

$N$  objects  $\rightarrow L$  clusters:

$$c(i) = l \Rightarrow i \in G_l \quad (1 \leq i \leq N, 1 \leq l \leq L)$$


$$Q(c, \{\mathbf{w}_l\}_1^L) = \sum_{l=1}^L \frac{W_l}{N_l^2} \sum_{c(i)=l} \sum_{c(j)=l} D_{ij}[\mathbf{w}_l]$$

- ▶ Each object  $i$  has  $n$  attributes:  $x_i = (x_{i1}, \dots, x_{in})$
- ▶  $\mathbf{w}_l = \{w_{kl}\}_1^n, 1 \leq l \leq L$ .
- ▶  $D_{ij}[\mathbf{w}_l] = \sum_{k=1}^n w_{kl} d_{ijk}$ , where  $d_{ijk} = \frac{|x_{ik} - x_{jk}|}{s_k}$ ;  
 $s_k = \frac{1}{N^2} \sum_{i=1}^N \sum_{j=1}^N |x_{ik} - x_{jk}|$ .
- ▶ Augmented to avoid clustering only on a single attribute:  
 $D_{ij}[\mathbf{w}_l] \rightarrow D_{ij}^{(\lambda)}[\mathbf{w}_l]$ .

# Minimization Scheme

- ▶ Start with an initial guess for  $w_{ki} = \frac{1}{n}$ ,  $\eta = \lambda$
- ▶ **outer WHILE**
- ▶ **inner WHILE**
- ▶ Fix  $w_{ki}$ , minimize the criterion w.r.t the encoder  $c(\cdot)$ 
  - ▶  $D_{ij}^{(\eta)}[\mathbf{w}] = -\eta \cdot \log \sum_{k=1}^n w_{ki} \cdot e^{-\frac{d_{ijk}}{\eta}}$ .  
 $d_{ijk} = \frac{|x_{ik} - x_{jk}|}{s_k}$ ,  $s_k = \frac{1}{N^2} \sum_{i=1}^N \sum_{j=1}^N |x_{ik} - x_{jk}|$ .
  - ▶  $D_{ij}^1[\mathbf{W}] = \max(D_{ij}^{(\eta)}[\mathbf{w}_{c(i|\mathbf{w})}], D_{ij}^{(\eta)}[\mathbf{w}_{c(j|\mathbf{w})}])$  (30)
  - ▶  $D_{ij}^{(2)}[\mathbf{W}] = \sum_{k=1}^n \max(w_{k,c(i|\mathbf{w})}, w_{k,c(j|\mathbf{w})}) d_{ijk}$  (33)
  - ▶  $KNN(i) = \{j | D_{ij} \leq d_{i(K)}\}$ .
- ▶ Fix  $c(\cdot)$ , minimize the criterion w.r.t the weights  $w_{ki}$ 
  - ▶  $w_{ki} = \frac{e^{-\frac{s_{ki}}{\lambda}}}{\sum_{k'=1}^n e^{-\frac{s_{ki}}{\lambda}}}$ ,  $S_{ki} = \frac{1}{K} \sum_{j \in KNN(i)} d_{ijk}$ .
- ▶ **END inner WHILE**
- ▶ Increment  $\eta = \eta + \alpha \cdot \lambda$
- ▶ **END outer WHILE**
- ▶ This iterative procedure is continued until convergence.

## Parallel COSA algorithm:




## Observed speedup on a real-world data

c22\_ceu.txt: (N=1000, n=13,000)

DELL R810 with 24 Xeon processors and a total 132GB memory

p	$p_x$	$p_y$	$p_z$	time
1	1	1	1	105417.38615 sec
2	1	1	2	51324.306341 sec
3	1	1	3	33812.867425 sec
4	1	1	4	26374.128593 sec
5	1	1	5	20340.902349 sec
6	1	1	6	17151.958595 sec
8	2	2	2	12327.523793 sec
12	2	2	3	9007.370069 sec
16	2	2	4	6559.531571 sec
20	2	2	5	5273.171368 sec
24	2	2	6	4778.485166 sec



# Conclusion

- ▶ Matching performance between MPI-COSA and Fortran-COSA executable (released by the original authors)
- ▶ MPI-COSA Source code was released publicly on [launchpad.net](https://launchpad.net)

Changes Files Help

~jgai/+junk/COSA : /parallel\_cosa2.c (revision 49)

browse files view revision view changes to this file download file

Line	Revision	Contents
1	13	#include <stdio.h>
2		#include <stdlib.h>
3	46	#include <assert.h>
4	13	#include <math.h>
5		#include <string.h>
6		#include <float.h>
7		#include <mpi.h>
8		#include <unistd.h>
9	36	#include "matrix_util.h"
10	19	#define GAI_DEBUG 0
11		
12	13	typedef struct{
13	36	int objectID;
14	13	float distance;
15		} NN_T; // nearest neighbor struct
16	36	
17	14	typedef struct
18		{
19		int p;
20		MPI_Comm comm;//entire grid
21	20	MPI_Comm slice_comm_along_z;// free_coords=[1,1,0]
22		MPI_Comm slice_comm_along_x;// free_coords=[0,1,1]
23	16	MPI_Comm depth_comm;// free_coords=[0,0,1]
24	18	MPI_Comm row_comm;// free_coords=[1,0,0]
25	25	MPI_Comm col_comm;// free_coords=[0,1,0]

# Performance Analysis

Sequential COSA:

$$T_1 = [4N^2 + (10n + 2K)N^2 + (5K + 5)nN + 3n + N] \cdot T^{flop} \cdot \mathbb{T}$$

Parallel COSA :

$$(p = p_x \cdot p_y \cdot p_z, \bar{N} = \frac{N}{p_x}, \bar{n} = \frac{n}{p_z})$$

$$T_q^{comp} = [5\bar{N}^2 + (10\bar{n} + 2K)\bar{N}^2 + (5K + 5)\bar{n}\bar{N} + (2K + 1)\bar{N} \cdot p_x \cdot K + 2\bar{N} + 2\bar{n}] \cdot T^{flop} \cdot \mathbb{T}$$

$$T_q^{comm} = [\bar{n} \cdot p_x p_y \cdot \log(p_x p_y) + \bar{N}^2 \cdot p_z \log(p_z) + 2\bar{N}^2 + 3\bar{N}K \cdot p_x \log(p_x) + \bar{n}\bar{N} \cdot p_x \log(p_x) + \bar{N} \cdot p_z \log(p_z) + p \log(p)] \cdot T^{Transmit} \cdot \mathbb{T}$$

$$\begin{aligned} T_q &= T_q^{comp} + T_q^{comm} \\ &= \left[ \frac{6N^2 n}{q} + \frac{N^2}{q} + \frac{N^2}{q} \log\left(\frac{N}{p}\right) + pK + \frac{nNK}{p} + \frac{n^2 N}{p} + \frac{2nN}{p} \right] \cdot T^{flop} \cdot \mathbb{T} \\ &\quad + [np \log(p) + 2pK] \cdot T^{Transmit} \cdot \mathbb{T} \end{aligned}$$



## Performance Analysis:

$$\begin{aligned}T_q &= T_q^{comp} + T_q^{comm} \\&= \left[ \frac{6N^2n}{q} + \frac{N^2}{q} + \frac{N^2}{q} \log\left(\frac{N}{p}\right) + pK + \frac{nNK}{p} + \frac{n^2N}{p} + \frac{2nN}{p} \right] \cdot T^{flop} \cdot \mathbb{T} \\&\quad + [np \log(p) + 2pK] \cdot T^{Transmit} \cdot \mathbb{T}\end{aligned}$$

- ▶  $T^{Transmit}$  is the time of sending a single message containing a floating point number.
- ▶  $T_q^{comm}$  is insignificant compared to  $T_q^{comp}$ , as  $N$  increases.  
For example

- ▶  $N = 11,000$ ,  $n = 2,000$ ,  $K = \sqrt{N} = 105$  and  $q = 16 (p = 4)$
- ▶  $T^{Transmit} = 10 \times 10^{-6}$  seconds,  $T^{flop} = 7 \times 10^{-11}$  seconds  
 $T_q^{comp} = 5.9875 \cdot \mathbb{T} \gg T_q^{comm} = 0.1684 \cdot \mathbb{T}$

Predicted Speedup:

$$S = \frac{T_1}{T_q} = \frac{[6N^2n + N^2 + N \log(N) + KnN + n^2N + 2nN] \cdot T^{flop} \cdot \mathbb{T}}{T_q^{comp} + T_q^{comm}}$$